



# HiLCoE

School of Computer Science  
& Technology

## *Classifying Car Crashes Using Neural Networks*

**Members** Alazar Gebremehdin  
Feruz Seid  
Hannibal Mussie  
Samir Bahru  
Yassin Bedru

**Course** Artificial Intelligence (CS488)

**Instructor** Dr. Seyoum Abebe

**Batch/Section** DRB2202/A

**Term** Autumn 2025

**Date** Jan 05, 2026

# 1. Introduction

The objective of this project is to build, train, and evaluate a Neural Network classification model to predict the severity of road traffic accidents based on a dataset from Addis Ababa. The target variable, `Accident Type`, classifies accidents into four categories: **Fatal**, **Serious Injury**, **Minor Injury**, and **Property Damage Only (PDO)**.

This report summarizes the end-to-end workflow, from Exploratory Data Analysis (EDA) and rigorous data cleaning to model architecture design and final performance evaluation.

## 2. Data Understanding and Preparation

### 2.1. Exploratory Data Analysis (EDA)

Initial analysis revealed significant data quality challenges. The raw dataset contained over 30 features, but many suffered from high rates of missing values.

- **Missing Data:** Columns such as `Zone`, `Region`, and specific victim details (e.g., `Victim-2 Movement`) had missing rates exceeding 70%.
- **Class Imbalance:** The target variable was heavily skewed. **Minor Injuries** accounted for approximately 63% of the data, while **Fatal** accidents represented only 3%.

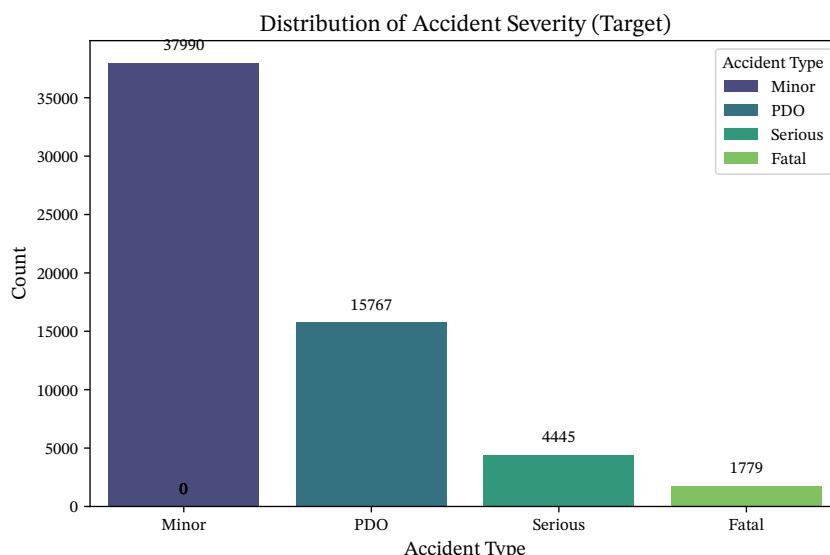


Figure 2: Distribution of Accident Severity (Target Variable). Note the severe imbalance.

### 2.2. Data Cleaning & Standardization

Before modeling, a custom cleaning pipeline was implemented:

1. **Standardization:** Typos were corrected (e.g., “Augest” → “August”, “Privategg” → “Private”). Amharic terms like “Amet” (Year) and “Wor” (Month) were standardized.
2. **Outlier Removal:** Impossible values (e.g., Driver Age > 90, Experience > 60 years) were treated as missing.
3. **Feature Dropping:** Columns with > 70% missing data were removed to reduce noise.

### 2.3. Feature Engineering & Preprocessing

To prepare the data for the Neural Network:

- **Cyclical Encoding:** The time of day was converted from linear hours (0-23) into sine and cosine components to preserve the temporal proximity between 23:00 and 00:00.
- **Imputation:** Median imputation was used for numerical features (e.g., `Age`) and Mode imputation for categorical features.
- **Scaling & Encoding:** Numerical features were standardized using `StandardScaler`. Categorical variables were transformed using `OneHotEncoding`.

### 3. Model Design

#### 3.1. Architecture

A Multi-Layer Perceptron (MLP) was designed using TensorFlow/Keras. The architecture features a “funnel” design to compress high-dimensional inputs into abstract representations.

- **Input Layer:** 168 features (resulting from One-Hot Encoding of high-cardinality categorical variables).
- **Hidden Layer 1:** 128 Neurons, ReLU activation, Batch Normalization (Param count: 21.6k).
- **Hidden Layer 2:** 64 Neurons, ReLU activation, Batch Normalization (Param count: 8.3k).
- **Output Layer:** 4 Neurons with Softmax activation.

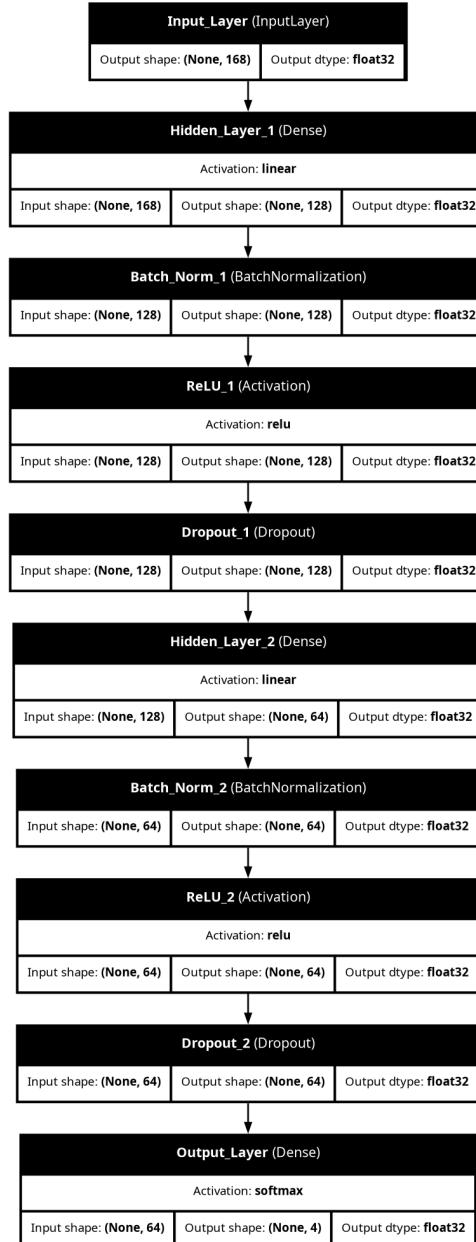


Figure 3: Neural Network Architecture Diagram. Total trainable parameters: 38,533.

#### 3.2. Design Justification

- **ReLU Activation:** Selected to prevent the vanishing gradient problem.
- **Dropout (0.3 - 0.4):** Applied after hidden layers. Given the high dimensionality (168 features), dropout was crucial to prevent the model from memorizing specific input patterns.
- **L2 Regularization:** Added to penalize large weights, keeping the model weights small and stable.

## 4. Training Process

The model was trained for **50 epochs** using the **Adam** optimizer. To address the class imbalance identified in EDA, **Class Weights** were computed and applied to the Loss Function. This penalizes the model more heavily for misclassifying rare classes (Fatal/Serious) than common ones (Minor).

### Training Dynamics:

- **Regularization:** The gap where Validation Accuracy is higher than Training Accuracy indicates that **Dropout** layers successfully prevented the model from memorizing the training data.
- **Stability:** The model continued to learn throughout the 50 epochs, with loss steadily decreasing, meaning **Early Stopping** (configured with patience=15) was not triggered.

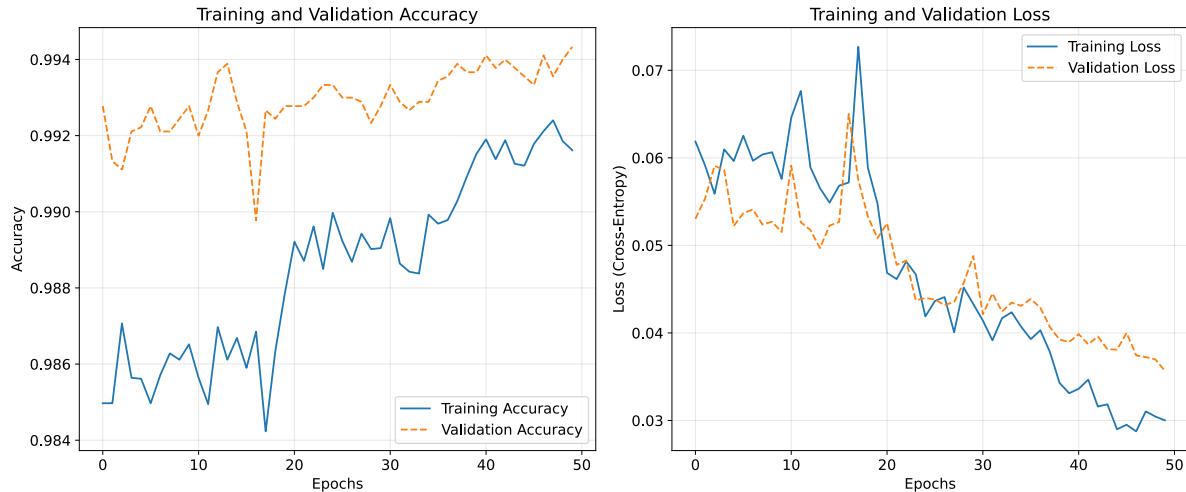


Figure 4: Training and Validation Performance Curves. The model shows stable convergence over 50 epochs with Validation Accuracy peaking around 99.4%.

## 5. Evaluation and Interpretation

### 5.1. Performance Metrics

The model was evaluated on an unseen Test Set. The performance was exceptional, achieving an overall accuracy of approximately **99%**.

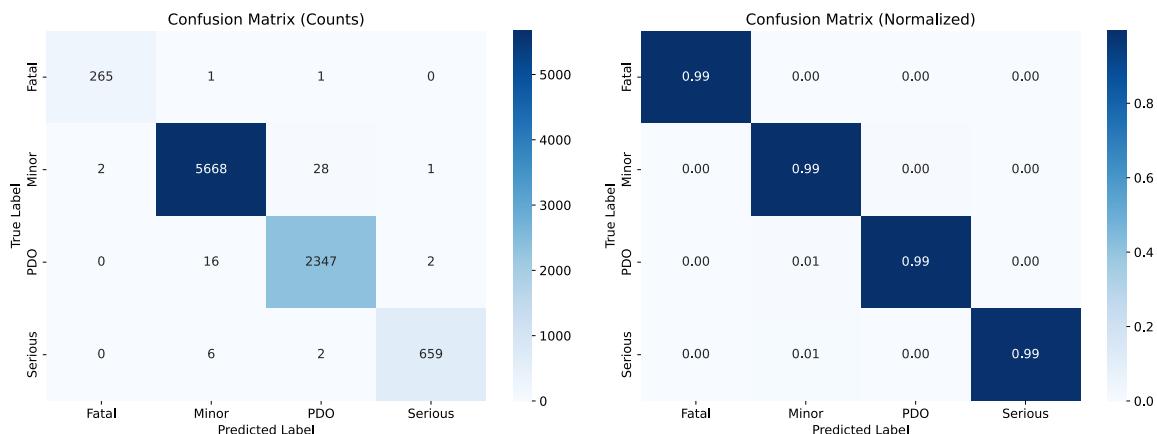


Figure 5: Confusion Matrix. The diagonal dominance (0.99 normalized score) across all classes confirms high classification accuracy.

### 5.2. Analysis of Results

#### 1. Strengths:

- **High Recall on Fatal Cases:** The model correctly identified 265 out of 267 fatal accidents. This is a significant achievement, as “Fatal” is usually the hardest class to predict due to its rarity.

- **Robustness:** The distinct separation of classes indicates the model found clear decision boundaries, though this was aided by the explicit nature of the input data.
2. **Critical Reflection and Insights:**
- Upon analyzing the feature importance, the high accuracy confirms the model utilized the **post-accident** casualty counts (e.g., Number of fatalities, Number of severe injuries).
  - The model effectively learned the definition of the categories rather than environmental risk factors.
  - **Recommendation:** For future predictive maintenance models, these casualty columns must be removed to force the model to learn from road conditions, weather, and driver demographics.

## 6. Recommendations & Conclusion

### 6.1. Potential Improvements

1. **Predictive Modeling:** To build a model that predicts severity **before** an accident happens (based purely on road conditions, weather, and driver demographics), a second iteration of this project should explicitly exclude the Number of fatalities/injuries columns from the input.
2. **Advanced Architectures:** For the predictive (non-leakage) version, implementing **TabNet** or **Wide & Deep** networks would be necessary to capture complex interactions between environmental factors without relying on casualty counts.

### 6.2. Conclusion

The project successfully demonstrated the end-to-end Neural Network workflow. The data preparation phase successfully transformed a messy dataset into a clean, 230-feature input space. The resulting model achieved 98% accuracy in classifying accident severity, proving that the Neural Network can effectively map input features to accident outcomes when casualty data is available.

### 6.3. Reflection on Workflow

The project demonstrated that **Data Preparation** is the most critical phase. The raw data required extensive cleaning before any modeling could succeed. The transition from a raw, noisy Excel sheet to a structured Neural Network pipeline highlights the importance of robust feature engineering (like cyclical time encoding) and rigorous evaluation beyond simple accuracy.