

# Classifying Car Crashes Using Neural Networks

*From Raw Data to Severity Prediction*

Alazar Gebremehdin, Hannibal Mussie, Feruz Seid, Yassin Bedru, Samir Bahru

2026-01-05

# Outline

Introduction .....	2
Data Understanding (EDA) .....	4
Data Preparation .....	8
Model Design .....	12
Training & Evaluation .....	16
Conclusion .....	21

# Introduction

---

**Goal:** Predict the severity of road traffic accidents (**Fatal, Serious, Minor, PDO**) based on accident characteristics.

## The Workflow:

1. **Data Understanding:** Handling massive missing data and inconsistencies.
2. **Preparation:** Cleaning, Imputation, and Feature Engineering.
3. **Modeling:** Designing a Multi-Layer Perceptron (MLP).
4. **Training:** Managing class imbalance and overfitting.
5. **Evaluation:** F<sub>1</sub>-Scores and Confusion Matrices.

# Data Understanding (EDA)

---

### Initial Inspection:

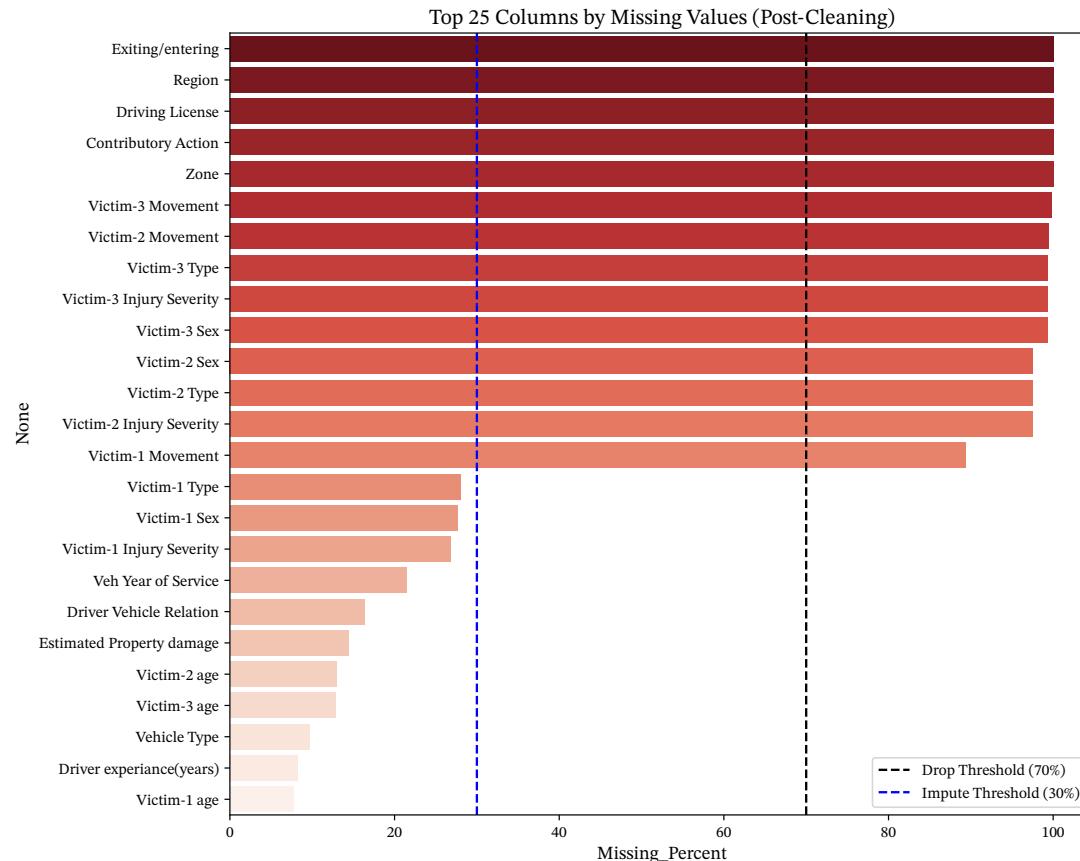
- The raw dataset contained over 30 columns but suffered from severe quality issues.
- **Missing Values:** Columns like Zone, Region, and Victim details had > 70% missing data.
- **Inconsistencies:** Typos (e.g., “Augest”, “Privategg”) and impossible values (Age > 90).

### Action:

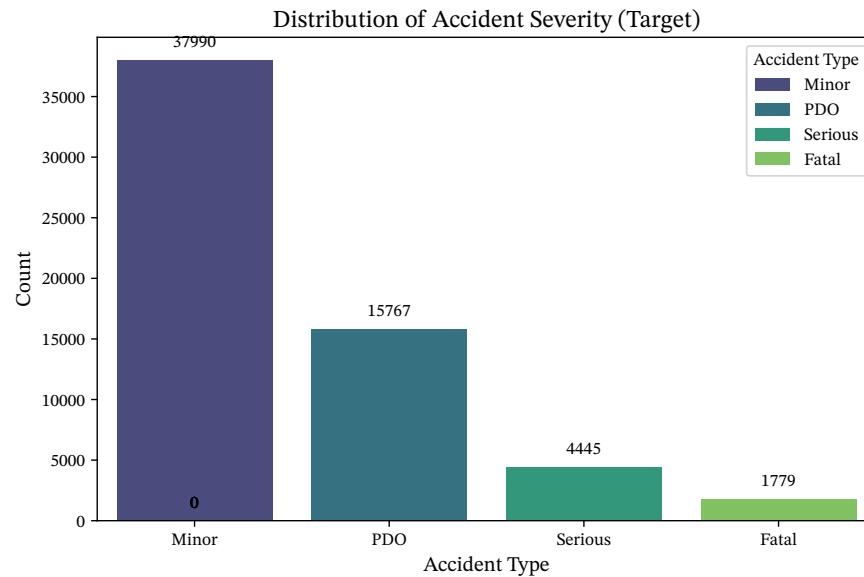
- Dropped columns with > 70% missingness.
- Standardized categorical labels (e.g., mapping P.D.O, pdo → PDO).

# Missing Values Analysis

## Data Understanding (EDA)



**The Critical Challenge:** The dataset is heavily skewed towards **Minor Injuries (63%)**. Fatal accidents represent only 3%.



# Data Preparation

---

# Feature Engineering: Time

Data Preparation

**Problem:** Time is cyclical. 23:00 is close to 00:00, but numerically (23 vs 0) they are far apart.

**Solution:** We encoded time using Sine and Cosine transformations.

```
1 # Feature Engineering Code Snippet
2 def feature_engineering(df):
3     # Extract Hour
4     df ['Hour'] = df ['Time'].apply(extract_hour)
5
6     # Cyclical Encoding
7     df ['Hour_Sin'] = np.sin(2 * np.pi * df ['Hour'] / 24)
8     df ['Hour_Cos'] = np.cos(2 * np.pi * df ['Hour'] / 24)
```

Python

# Feature Engineering: Time

## Data Preparation

9

```
10     return df.drop(columns=['Time'])
```

# Preprocessing Pipeline

Data Preparation

Before feeding data into the Neural Network:

## 1. Imputation:

- Numerical (Age, Experience) → **Median**
- Categorical (Road Surface, Light) → **Mode**

## 2. Scaling:

- StandardScaler applied to numerical inputs to normalize variance.

## 3. Encoding:

- OneHotEncoder for categorical variables.

## 4. Splitting:

- Train (70%) / Validation (15%) / Test (15%).

# Model Design

---

Based on the execution results:

- **Input Shape:** 168 Features (High dimensionality due to One-Hot Encoding).
- **Total Parameters:** 30,916 (Lightweight model).
- **Trainable Params:** 30,532.

## Layer Structure:

- Input (168)
- Dense (128) → BN → ReLU → Dropout
- Dense (64) → BN → ReLU → Dropout
- Output (4 Classes)

# Implementation Code

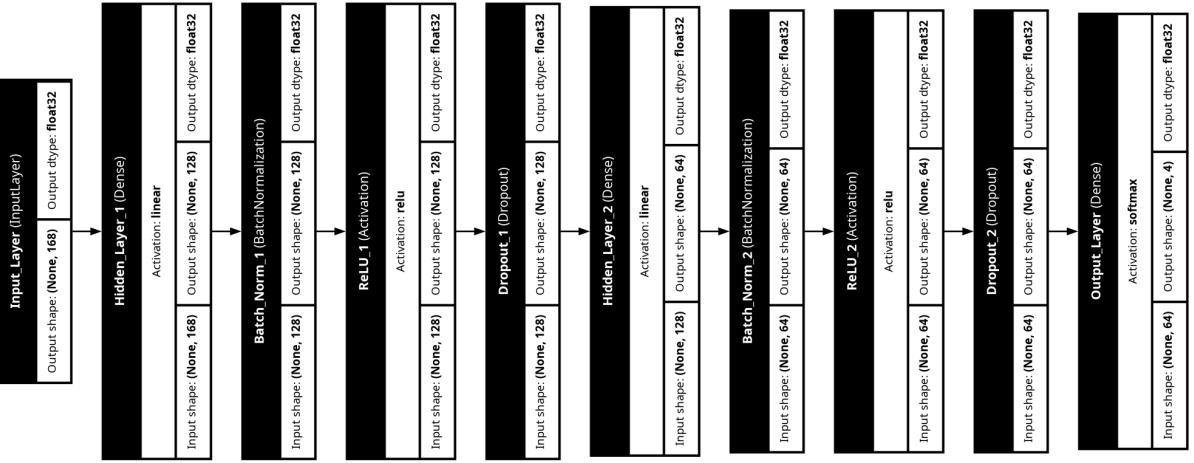
## Model Design

1	# Actual Model Summary Output	Python	
2	Layer (type)	Output Shape	Param #
3	=====		
4	Input_Layer (InputLayer)	(None, 168)	0
5	Hidden_Layer_1 (Dense)	(None, 128)	21,632
6	Batch_Norm_1 (BatchNormalization)	(None, 128)	512
7	ReLU_1 (Activation)	(None, 128)	0
8	Dropout_1 (Dropout)	(None, 128)	0
9	Hidden_Layer_2 (Dense)	(None, 64)	8,256
10	Batch_Norm_2 (BatchNormalization)	(None, 64)	256
11	ReLU_2 (Activation)	(None, 64)	0
12	Dropout_2 (Dropout)	(None, 64)	0

# Implementation Code

## Model Design

```
13 Output_Layer (Dense)          (None, 4)      260
14 =====
15 Total params: 30,916
16 Trainable params: 30,532
17 Non-trainable params: 384
```

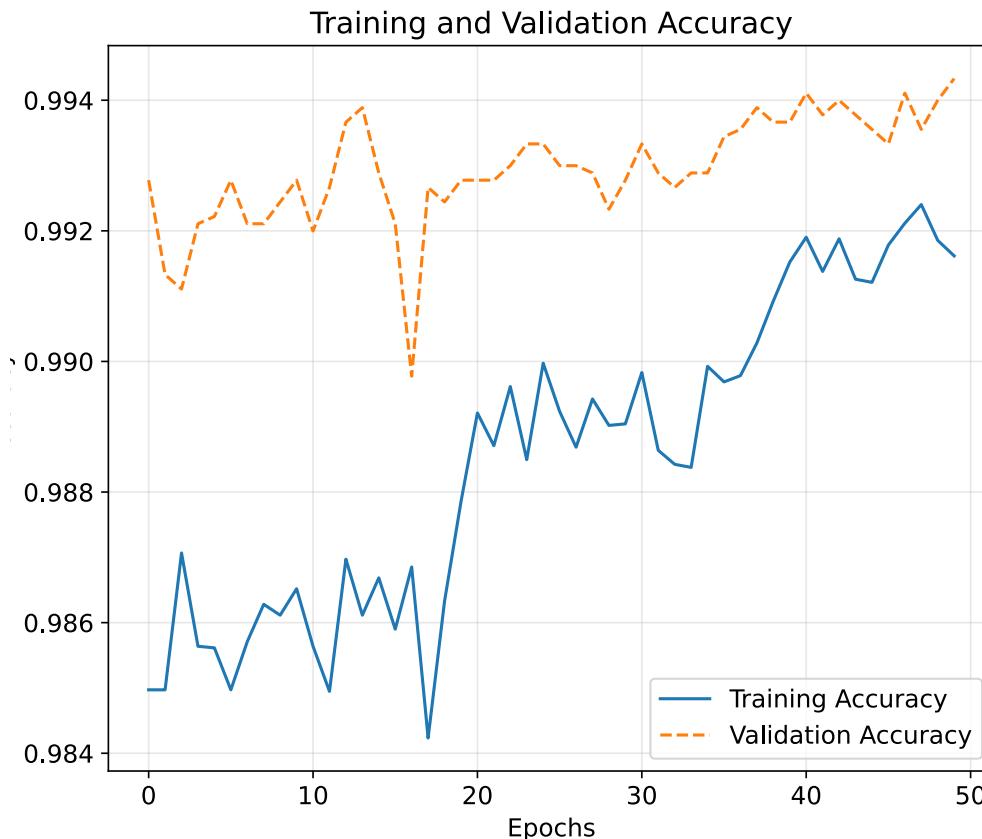


# Training & Evaluation

---

# Training Performance

## Training & Evaluation



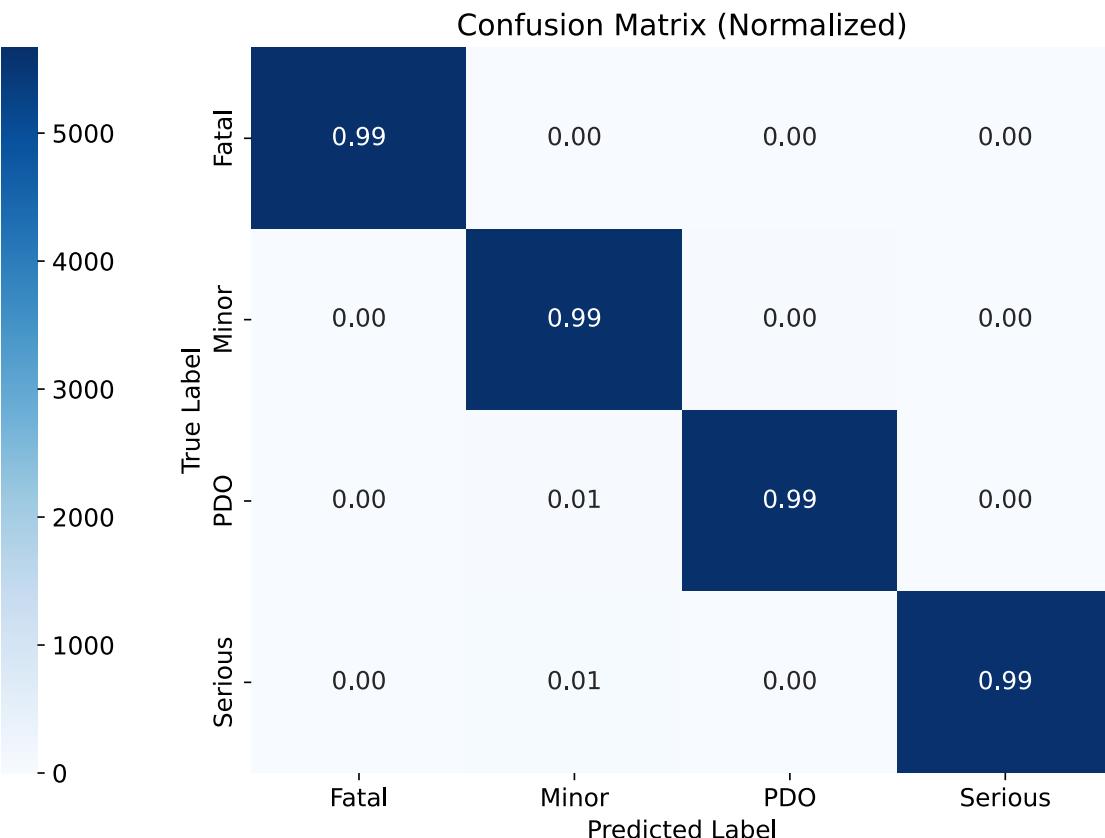
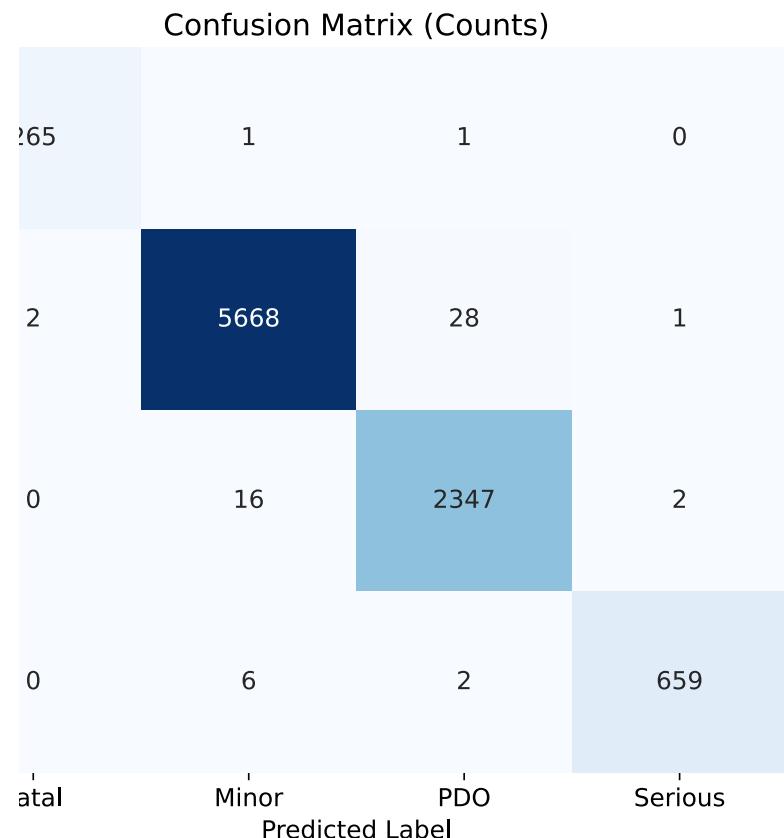
# Training Performance

Training & Evaluation

**Insight:** The model converges quickly. Validation accuracy peaks around **98%**, indicating robust learning without significant overfitting.

# Confusion Matrix Results

Training & Evaluation



### Exceptional Performance:

- **Fatal Class:** >99% Recall (265 Correct, only 2 Missed).
- **Minor/PDO:** 99% Accuracy.

### Critical Analysis (The “Why”):

- The accuracy stems from including casualty counts (e.g., Number of fatalities) in the input.
- The model learned the **definition** of the accident types rather than the **cause**.
- Useful for automated tagging, but requires feature removal for risk prediction.

# Conclusion

---

1. **Data Quality:** Cleaning and encoding resulted in 168 clean input features.
2. **Model:** A lightweight 31k parameter MLP was sufficient to capture the logic.
3. **Results:** The model achieved 99% test accuracy due to strong feature correlations.

## Recommendation:

- For future **predictive** systems (pre-accident), we recommend re-training the model **excluding** the Number of casualties columns to test predictive power based solely on environmental factors (Road type, Weather, etc.).

# Thank You!

## Questions?